# Supplementary Table S1

Per-label metrics for all main systems and four fixed LLM configurations

| System | Label | Precision | Recall | F1 | Accuracy | ROC-AUC | AP |
|---|---|---|---|---|---|---|---|
| All-positive | Toxic | 0.354 | 1.000 | 0.523 | 0.354 | — | — |
| | Offensive | 0.233 | 1.000 | 0.377 | 0.233 | — | — |
| | Hate | 0.040 | 1.000 | 0.075 | 0.040 | — | — |
| LR-TFIDF | Toxic | 0.408 | 0.786 | 0.535 | 0.520 | 0.630 | 0.506 |
| | Offensive | 0.246 | 0.547 | 0.333 | 0.540 | 0.609 | 0.332 |
| | Hate | 0.072 | 0.405 | 0.117 | 0.719 | 0.582 | 0.075 |
| SVM-TFIDF | Toxic | 0.399 | 0.760 | 0.522 | 0.507 | 0.618 | 0.491 |
| | Offensive | 0.240 | 0.587 | 0.331 | 0.500 | 0.613 | 0.353 |
| | Hate | 0.064 | 0.310 | 0.100 | 0.723 | 0.588 | 0.076 |
| LR-Count | Toxic | 0.390 | 0.705 | 0.499 | 0.502 | 0.583 | 0.449 |
| | Offensive | 0.260 | 0.618 | 0.356 | 0.490 | 0.575 | 0.315 |
| | Hate | 0.044 | 0.305 | 0.075 | 0.711 | 0.567 | 0.074 |
| XLM-R | Toxic | 0.442 | 0.802 | 0.560 | 0.549 | 0.659 | 0.514 |
| | Offensive | 0.326 | 0.733 | 0.435 | 0.554 | 0.690 | 0.413 |
| | Hate | 0.031 | 0.370 | 0.054 | 0.648 | 0.498 | 0.087 |
| mBERT | Toxic | 0.402 | 0.817 | 0.533 | 0.493 | 0.633 | 0.525 |
| | Offensive | 0.341 | 0.637 | 0.414 | 0.611 | 0.689 | 0.433 |
| | Hate | 0.045 | 0.367 | 0.075 | 0.640 | 0.522 | 0.107 |
| LLM zero-shot binary | Toxic | 0.739 | 0.714 | 0.725 | 0.809 | — | — |
| | Offensive | 0.705 | 0.599 | 0.645 | 0.847 | — | — |
| | Hate | 0.450 | 0.324 | 0.370 | 0.957 | — | — |
| LLM zero-shot score | Toxic | 0.661 | 0.771 | 0.708 | 0.776 | 0.845 | 0.750 |
| | Offensive | 0.665 | 0.660 | 0.642 | 0.828 | 0.848 | 0.711 |
| | Hate | 0.163 | 0.333 | 0.141 | 0.771 | 0.648 | 0.308 |
| LLM six-shot binary | Toxic | 0.737 | 0.684 | 0.709 | 0.801 | — | — |
| | Offensive | 0.657 | 0.675 | 0.658 | 0.838 | — | — |
| | Hate | 0.424 | 0.357 | 0.377 | 0.953 | — | — |
| LLM six-shot score | Toxic | 0.764 | 0.750 | 0.755 | 0.829 | 0.830 | 0.726 |
| | Offensive | 0.712 | 0.647 | 0.664 | 0.848 | 0.819 | 0.621 |
| | Hate | 0.446 | 0.433 | 0.326 | 0.784 | 0.636 | 0.248 |
| LLM (selected) | Toxic | 0.743 | 0.781 | 0.757 | 0.821 | 0.832 | 0.734 |
| | Offensive | 0.676 | 0.655 | 0.641 | 0.828 | 0.822 | 0.641 |
| | Hate | 0.446 | 0.433 | 0.326 | 0.784 | 0.634 | 0.263 |

Scores are dimensionless means over five test folds. Transformer metrics are first averaged over five seeds within each fold. AP: average precision. Dashes denote ranking metrics not reported for constant-reference or binary-only output under the prespecified protocol. LLM (selected) uses validation-only configuration selection per fold. Standard deviations are provided in label_metrics_full.csv. These are preliminary results on 554 comments.